

```
pip install pandas numpy scikit-learn tensorflow
```

→ Requirement already satisfied: pandas in /usr/local/lib/python3.11/dist-packages (2.2.2)  
Requirement already satisfied: numpy in /usr/local/lib/python3.11/dist-packages (2.0.2)  
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.11/dist-packages (1.6.1)  
Requirement already satisfied: tensorflow in /usr/local/lib/python3.11/dist-packages (2.18.0)  
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.11/dist-packages (from pandas) (2.8.2)  
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-packages (from pandas) (2025.2)  
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-packages (from pandas) (2025.2)  
Requirement already satisfied: scipy>=1.6.0 in /usr/local/lib/python3.11/dist-packages (from scikit-learn) (1.14.1)  
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.11/dist-packages (from scikit-learn) (1.4.2)  
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.11/dist-packages (from scikit-learn) (3.6.0)  
Requirement already satisfied: absl-py>=1.0.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.4.0)  
Requirement already satisfied: astunparse>=1.6.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.6.3)  
Requirement already satisfied: flatbuffers>=24.3.25 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (25.2.10)  
Requirement already satisfied: gast!=0.5.0,!0.5.1,!0.5.2,>0.2.1 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (0.6)  
Requirement already satisfied: google-pasta>=0.1.1 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (0.2.0)  
Requirement already satisfied: libclang>=13.0.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (18.1.1)  
Requirement already satisfied: opt-einsum>=2.3.2 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (3.4.0)  
Requirement already satisfied: packaging in /usr/local/lib/python3.11/dist-packages (from tensorflow) (24.2)  
Requirement already satisfied: protobuf!=4.21.0,!4.21.1,!4.21.2,!4.21.3,!4.21.4,!4.21.5,<6.0.0dev,>=3.20.3 in /usr/local/lib/p  
Requirement already satisfied: requests<3,>=2.21.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (2.32.3)  
Requirement already satisfied: setuptools in /usr/local/lib/python3.11/dist-packages (from tensorflow) (75.2.0)  
Requirement already satisfied: six>=1.12.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.17.0)  
Requirement already satisfied: termcolor>=1.1.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (3.0.1)  
Requirement already satisfied: typing-extensions>=3.6.6 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (4.13.2)  
Requirement already satisfied: wrapt>=1.11.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.17.2)  
Requirement already satisfied: grpcio<2.0,>=1.24.3 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.71.0)  
Requirement already satisfied: tensorboard<2.19,>=2.18 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (2.18.0)  
Requirement already satisfied: keras>=3.5.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (3.8.0)  
Requirement already satisfied: h5py>=3.11.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (3.13.0)  
Requirement already satisfied: ml-dtypes<0.5.0,>=0.4.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (0.4.1)  
Requirement already satisfied: tensorflow-io-gcs-filesystem>=0.23.1 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (0  
Requirement already satisfied: wheel<1.0,>=0.23.0 in /usr/local/lib/python3.11/dist-packages (from astunparse>=1.6.0->tensorflow) (0  
Requirement already satisfied: rich in /usr/local/lib/python3.11/dist-packages (from keras>=3.5.0->tensorflow) (13.9.4)  
Requirement already satisfied: namex in /usr/local/lib/python3.11/dist-packages (from keras>=3.5.0->tensorflow) (0.0.8)  
Requirement already satisfied: optree in /usr/local/lib/python3.11/dist-packages (from keras>=3.5.0->tensorflow) (0.15.0)  
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0->tens  
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0->tensorflow) (3.10  
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0->tens  
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0->tens  
Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.11/dist-packages (from tensorboard<2.19,>=2.18->tens  
Requirement already satisfied: tensorboard-data-server<0.8.0,>=0.7.0 in /usr/local/lib/python3.11/dist-packages (from tensorboard<2  
Requirement already satisfied: werkzeug>=1.0.1 in /usr/local/lib/python3.11/dist-packages (from tensorboard<2.19,>=2.18->tens  
Requirement already satisfied: MarkupSafe>=2.1.1 in /usr/local/lib/python3.11/dist-packages (from werkzeug>=1.0.1->tens  
Requirement already satisfied: markdown-it-py>=2.2.0 in /usr/local/lib/python3.11/dist-packages (from rich->keras>=3.5.0->tens  
Requirement already satisfied: pygments<3.0.0,>=2.13.0 in /usr/local/lib/python3.11/dist-packages (from rich->keras>=3.5.0->tens  
Requirement already satisfied: mdurl~0.1 in /usr/local/lib/python3.11/dist-packages (from markdown-it-py>=2.2.0->rich->keras>=3.5.0

```
import pandas as pd
import numpy as np
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences

from google.colab import files
uploaded = files.upload()
```

Choose File No file chosen

```
import os
from google.colab import drive
drive.mount('/content/drive')
train_directory = '/content/drive/MyDrive/EmotionsDataset_Splitted/data/test'
files = os.listdir(train_directory)
print(files)
```

→ Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force\_remount=True).  
['sad', 'nothing', 'angry', 'happy']

```
import chardet

def read_file_safely(file_path):
    with open(file_path, 'rb') as f:
        raw_data = f.read()
        result = chardet.detect(raw_data)
        encoding = result['encoding']
```

```
with open(file_path, 'r', encoding=encoding, errors='ignore') as f:
    return f.read().strip()

def load_data_from_folder(folder_path):
    texts = []
    labels = []

    for label in os.listdir(folder_path):
        label_folder = os.path.join(folder_path, label)
        if os.path.isdir(label_folder):
            for file in os.listdir(label_folder):
                file_path = os.path.join(label_folder, file)
                try:
                    text = read_file_safely(file_path)
                    texts.append(text)
                    labels.append(label)
                except Exception as e:
                    print(f"Skipping file {file_path} due to error: {e}")

    return texts, labels

test_folder = '/content/drive/MyDrive/EmotionsDataset_Splitted/data/test'
texts, labels = load_data_from_folder(test_folder)
```

Next steps: [Generate code with test\\_df](#)

```
→ (788, 2)
emotion
nothing    219
angry      213
happy      207
sad        149
Name: count, dtype: int64
```

```
CONFIG = {
    "BATCH_SIZE": 32,
    "IM_SIZE": 128,
    "LEARNING_RATE": 1e-3,
    "N_EPOCHS": 5,
    "DROPOUT_RATE": 0.3,
    "REGULARIZATION_RATE": 0.001,
    "NUM_CLASSES": 3,
    "CLASS_NAMES": ["angry", "happy", "nothing", "sad"]}
```

```
import tensorflow as tf
train_df = tf.keras.utils.image_dataset_from_directory(
    train_directory,
    labels='inferred',
    label_mode='int',
    class_names=CONFIG["CLASS_NAMES"],
    color_mode='rgb',
    batch_size=CONFIG["BATCH_SIZE"],
    image_size=(CONFIG["IM_SIZE"], CONFIG["IM_SIZE"]))
```

```
shuffle=True,
seed=42
)
```

→ Found 788 files belonging to 4 classes.

```
from tensorflow.keras.layers import Resizing, Rescaling
resize_rescale_layers = tf.keras.Sequential([
    Resizing(CONFIG['IM_SIZE'], CONFIG['IM_SIZE']),
    Rescaling(1./255)
])
train_df = train_df.map(lambda x, y: (resize_rescale_layers(x), y), num_parallel_calls=tf.data.AUTOTUNE)
train_df = train_df.prefetch(buffer_size=tf.data.AUTOTUNE)

import matplotlib.pyplot as plt
plt.figure(figsize=(12, 12))
for images, labels in train_df.take(1):
    for i in range(16):
        ax = plt.subplot(4, 4, i + 1)
        plt.imshow(images[i].numpy())
        plt.title(CONFIG["CLASS_NAMES"][labels[i].numpy()])
        plt.axis("off")
    plt.show()
```

→

nothing



nothing



happy



sad



happy



happy



sad



nothing



nothing



nothing



sad



happy



nothing



happy



nothing



nothing



```
plt.figure(figsize = (12,12))
for images, labels in train_df.take (1) :
```

```
for i in range (16):
    ax = plt.subplot(4,4, i+1)
    plt. imshow(images [i]. numpy())
    plt.title(CONFIG["CLASS_NAMES"][labels[i]])

    plt.axis ("off")
```



```
for i in train_df. take (1) :
    print(i)
```



```
[14./0b49540e-01, 4.00/89240e-01, 4.200/0480e-01],
[2.62132376e-01, 2.96139717e-01, 2.59252459e-01],
[1.04718141e-01, 1.40870109e-01, 1.30575985e-01],
...,
[6.73897088e-01, 5.38725495e-01, 4.30759817e-01],
[6.62316203e-01, 5.21139741e-01, 4.19178933e-01],
[6.45098090e-01, 5.03921628e-01, 4.01960820e-01]],

...,

[[1.67218149e-01, 4.43995118e-01, 5.29840708e-01],
[1.60477951e-01, 4.19424057e-01, 5.82107902e-01],
[3.39460820e-02, 1.87806383e-01, 3.14460814e-01],
...,
[2.38357857e-02, 3.19791675e-01, 3.50306392e-01],
[9.03186351e-02, 3.93995136e-01, 4.32659328e-01],
[5.58210835e-02, 3.52879912e-01, 4.01899546e-01]],

[[6.98529482e-02, 2.88296580e-01, 3.69852960e-01],
[2.66544130e-02, 1.95526972e-01, 3.47058833e-01],
[1.22303925e-01, 3.27205896e-01, 4.30698544e-01],
...,
[5.04289232e-02, 3.71691197e-01, 3.95772070e-01],
[4.18504924e-02, 3.07781875e-01, 3.37009817e-01],
[2.57965699e-02, 3.13419133e-01, 3.38848054e-01]],

[[1.40931383e-01, 3.35784346e-01, 4.13235307e-01],
[7.29779452e-02, 2.35477954e-01, 3.25612754e-01],
[1.84926480e-01, 4.21629936e-01, 4.80882376e-01],
...,
[6.55024573e-02, 3.43382359e-01, 3.95588249e-01],
[6.25000000e-02, 3.57536793e-01, 3.96446109e-01],
[7.24877492e-02, 4.18504924e-01, 4.40686315e-01]]],  

dtype=float32), <tf.Tensor: shape=(32,), dtype=int32, numpy=  

array([2, 0, 0, 2, 2, 1, 2, 0, 3, 0, 3, 1, 2, 1, 0, 2, 3, 0, 2, 0, 3, 1,  

3, 2, 3, 1, 1, 2, 3, 0, 2, 2], dtype=int32)>)
```

## ▼ CNN MODEL

```
import os
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.preprocessing.image import ImageDataGenerator
import matplotlib.pyplot as plt

import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Input, Conv2D, MaxPooling2D, Flatten, Dense, Dropout
CONFIG = {
    'IM_SIZE': 128,
    'DROPOUT_RATE': 0.5,
    'NUM_CLASSES': 10,
    'LEARNING_RATE': 0.001,
    'N_EPOCHS': 20
}
def create_cnn_model():
    model = Sequential([
        Input(shape=(CONFIG['IM_SIZE'], CONFIG['IM_SIZE'], 3)),
        Conv2D(32, (3, 3), activation='relu', padding='same'),
        MaxPooling2D((2, 2)),
        Conv2D(64, (3, 3), activation='relu', padding='same'),
        MaxPooling2D((2, 2)),
        Conv2D(128, (3, 3), activation='relu', padding='same'),
        MaxPooling2D((2, 2)),
        Flatten(),
        Dense(128, activation='relu'),
        Dropout(CONFIG['DROPOUT_RATE']),
        Dense(CONFIG['NUM_CLASSES'], activation='softmax')
    ])
    return model
cnn_model = create_cnn_model()
cnn_model.compile(
    optimizer=tf.keras.optimizers.Adam(learning_rate=CONFIG['LEARNING_RATE']),
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy']
)
cnn_model.summary()
```

## Model: "sequential\_3"

Layer (type)	Output Shape	Param #
conv2d_10 (Conv2D)	(None, 128, 128, 32)	896
max_pooling2d_10 (MaxPooling2D)	(None, 64, 64, 32)	0
conv2d_11 (Conv2D)	(None, 64, 64, 64)	18,496
max_pooling2d_11 (MaxPooling2D)	(None, 32, 32, 64)	0
conv2d_12 (Conv2D)	(None, 32, 32, 128)	73,856
max_pooling2d_12 (MaxPooling2D)	(None, 16, 16, 128)	0
flatten_2 (Flatten)	(None, 32768)	0
dense_4 (Dense)	(None, 128)	4,194,432
dropout_2 (Dropout)	(None, 128)	0
dense_5 (Dense)	(None, 10)	1,290

Total params: 4,288,970 (16.36 MB)

```

from tensorflow.keras.losses import SparseCategoricalCrossentropy
metrics = [
    tf.keras.metrics.SparseCategoricalAccuracy(name='accuracy'),
    tf.keras.metrics.SparseTopKCategoricalAccuracy(k=2, name='top_k_accuracy')
]
loss_function = SparseCategoricalCrossentropy()
cnn_model.compile(
    optimizer=tf.keras.optimizers.Adam(learning_rate=CONFIG['LEARNING_RATE']),
    loss=loss_function,
    metrics=metrics
)
history_cnn = cnn_model.fit(
    train_df,
    epochs=CONFIG['N_EPOCHS'],
    verbose=1
)

→ Epoch 1/20
25/25 ━━━━━━━━━━ 35s 1s/step - accuracy: 0.3968 - loss: 1.5981 - top_k_accuracy: 0.6352
Epoch 2/20
25/25 ━━━━━━━━━━ 38s 1s/step - accuracy: 0.5320 - loss: 1.0160 - top_k_accuracy: 0.7999
Epoch 3/20
25/25 ━━━━━━━━━━ 40s 1s/step - accuracy: 0.5743 - loss: 0.8940 - top_k_accuracy: 0.7796
Epoch 4/20
25/25 ━━━━━━━━━━ 30s 1s/step - accuracy: 0.5424 - loss: 0.8895 - top_k_accuracy: 0.7993
Epoch 5/20
25/25 ━━━━━━━━━━ 40s 1s/step - accuracy: 0.6129 - loss: 0.7696 - top_k_accuracy: 0.8425
Epoch 6/20
25/25 ━━━━━━━━━━ 28s 1s/step - accuracy: 0.6512 - loss: 0.7511 - top_k_accuracy: 0.8580
Epoch 7/20
25/25 ━━━━━━━━━━ 41s 1s/step - accuracy: 0.6633 - loss: 0.7107 - top_k_accuracy: 0.8791
Epoch 8/20
25/25 ━━━━━━━━━━ 28s 1s/step - accuracy: 0.7131 - loss: 0.6619 - top_k_accuracy: 0.8981
Epoch 9/20
25/25 ━━━━━━━━━━ 28s 1s/step - accuracy: 0.6947 - loss: 0.6415 - top_k_accuracy: 0.8858
Epoch 10/20
25/25 ━━━━━━━━━━ 41s 1s/step - accuracy: 0.7456 - loss: 0.5920 - top_k_accuracy: 0.8972
Epoch 11/20
25/25 ━━━━━━━━━━ 28s 1s/step - accuracy: 0.7544 - loss: 0.5614 - top_k_accuracy: 0.9377
Epoch 12/20
25/25 ━━━━━━━━━━ 41s 1s/step - accuracy: 0.7985 - loss: 0.5052 - top_k_accuracy: 0.9367
Epoch 13/20
25/25 ━━━━━━━━━━ 41s 1s/step - accuracy: 0.7958 - loss: 0.5455 - top_k_accuracy: 0.9312
Epoch 14/20
25/25 ━━━━━━━━━━ 29s 1s/step - accuracy: 0.7979 - loss: 0.4614 - top_k_accuracy: 0.9476
Epoch 15/20
25/25 ━━━━━━━━━━ 41s 1s/step - accuracy: 0.8473 - loss: 0.4237 - top_k_accuracy: 0.9541
Epoch 16/20
25/25 ━━━━━━━━━━ 41s 1s/step - accuracy: 0.8564 - loss: 0.3843 - top_k_accuracy: 0.9430
Epoch 17/20
25/25 ━━━━━━━━━━ 29s 1s/step - accuracy: 0.8751 - loss: 0.3255 - top_k_accuracy: 0.9730
Epoch 18/20
25/25 ━━━━━━━━━━ 28s 1s/step - accuracy: 0.8525 - loss: 0.3466 - top_k_accuracy: 0.9677
Epoch 19/20
25/25 ━━━━━━━━━━ 29s 1s/step - accuracy: 0.8870 - loss: 0.2928 - top_k_accuracy: 0.9792
Epoch 20/20
25/25 ━━━━━━━━━━ 41s 1s/step - accuracy: 0.9108 - loss: 0.2620 - top_k_accuracy: 0.9820

```

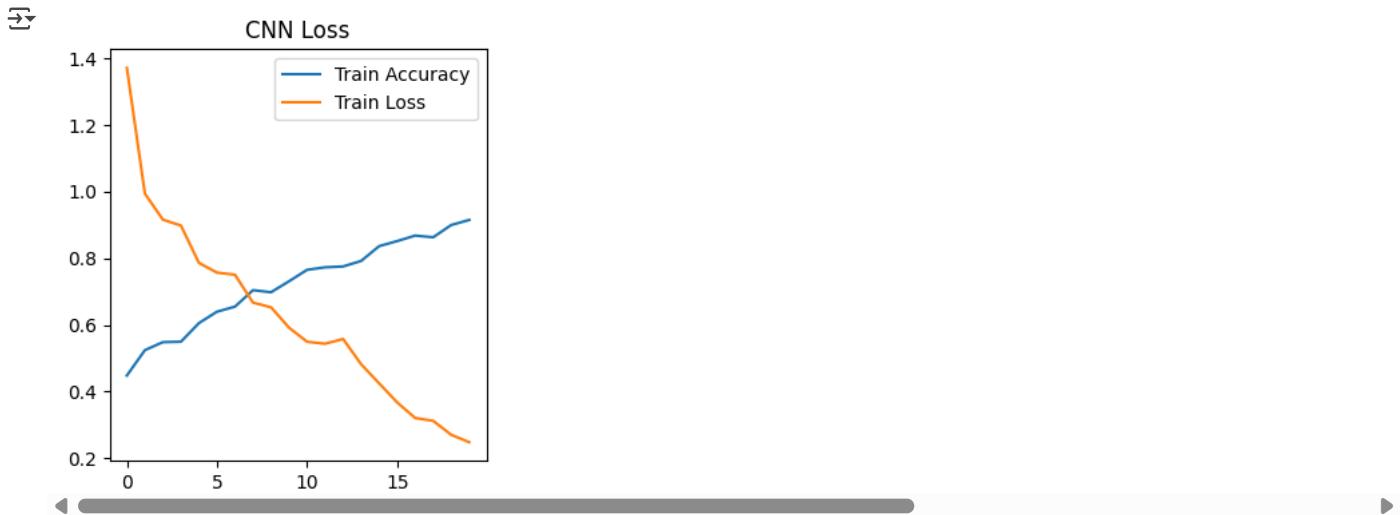
```

plt.figure(figsize=(8, 4))
plt.subplot(1, 2, 1)

```

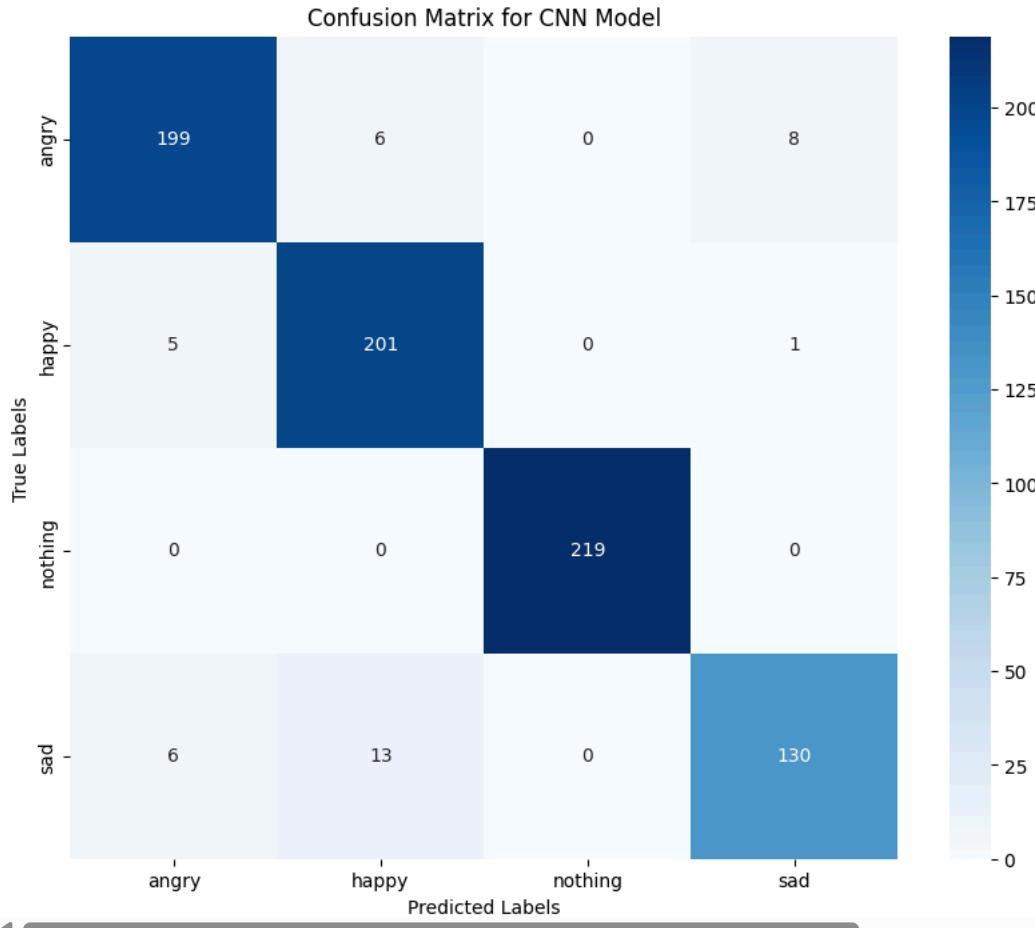
```
plt.plot(history_cnn.history['accuracy'], label='Train Accuracy')
plt.plot(history_cnn.history['loss'], label='Train Loss')
plt.legend()
plt.title('CNN Accuracy')
plt.title('CNN Loss')

plt.show()
```



```
import tensorflow as tf
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix
import seaborn as sns
import numpy as np
test_dataset = tf.keras.utils.image_dataset_from_directory(
    '/content/drive/MyDrive/EmotionsDataset_Splitted/data/test',
    image_size=(CONFIG['IM_SIZE'], CONFIG['IM_SIZE']),
    batch_size=32
)
CONFIG['CLASS_NAMES'] = test_dataset.class_names
CNN_INPUT_SIZE = cnn_model.input_shape[1:3]
true_labels = []
predicted_labels = []
for images, labels in test_dataset:
    images = images / 255.0
    preds = cnn_model(images, training=False)
    pred_label_idx = tf.argmax(preds, axis=1).numpy()
    true_labels.extend(labels.numpy())
    predicted_labels.extend(pred_label_idx)
true_labels = np.array(true_labels)
predicted_labels = np.array(predicted_labels)
cm = confusion_matrix(true_labels, predicted_labels)
plt.figure(figsize=(10, 8))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=CONFIG['CLASS_NAMES'], yticklabels=CONFIG['CLASS_NAMES'])
plt.title("Confusion Matrix for CNN Model")
plt.xlabel("Predicted Labels")
plt.ylabel("True Labels")
plt.show()
```

→ Found 788 files belonging to 4 classes.



```
cnn_eval_results = cnn_model.evaluate(train_df, verbose=1)
print(f"\nCNN Model Evaluation Results:\nLoss: {cnn_eval_results[0]:.4f}, Accuracy: {cnn_eval_results[1]:.4f}, Top-2 Accuracy: {cnn_eval_results[2]:.4f}")
```

→ 25/25 ————— 10s 397ms/step - accuracy: 0.9374 - loss: 0.1537 - top\_k\_accuracy: 0.9939

```
CNN Model Evaluation Results:
Loss: 0.1359, Accuracy: 0.9505, Top-2 Accuracy: 0.9949
```

```
import cv2
import matplotlib.pyplot as plt

img_path = '/content/drive/MyDrive/EmotionsDataset_Splitted/data/test/sad/sad_140.jpg'
test_image = cv2.imread(img_path)

if test_image is None:
    print("Image not found or failed to load:", img_path)
else:
    test_image = cv2.cvtColor(test_image, cv2.COLOR_BGR2RGB)
    plt.imshow(test_image)
    plt.axis('off')
    plt.show()
```



```
import cv2
import matplotlib.pyplot as plt

img_path = '/content/drive/MyDrive/EmotionsDataset_Splitted/data/test/happy/happy_140.jpg'
test_image = cv2.imread(img_path)

if test_image is None:
    print("Image not found or failed to load:", img_path)
else:
    test_image = cv2.cvtColor(test_image, cv2.COLOR_BGR2RGB)
    plt.imshow(test_image)
    plt.axis('off')
    plt.show()
```



```
import cv2
import matplotlib.pyplot as plt

img_path = '/content/drive/MyDrive/EmotionsDataset_Splitted/data/test/angry/angry_140.jpg'
test_image = cv2.imread(img_path)

if test_image is None:
    print("Image not found or failed to load:", img_path)
else:
    test_image = cv2.cvtColor(test_image, cv2.COLOR_BGR2RGB)
    plt.imshow(test_image)
    plt.axis('off')
    plt.show()
```



```
import cv2
import matplotlib.pyplot as plt

img_path = '/content/drive/MyDrive/EmotionsDataset_Splitted/data/test/nothing/nothing_140.jpg'
test_image = cv2.imread(img_path)

if test_image is None:
    print("Image not found or failed to load:", img_path)
else:
    test_image = cv2.cvtColor(test_image, cv2.COLOR_BGR2RGB)
    plt.imshow(test_image)
    plt.axis('off')
    plt.show()
```



```
CONFIG = {
    'IM_SIZE': 128,
    'DROPOUT_RATE': 0.5,
    'NUM_CLASSES': 3,
    'LEARNING_RATE': 0.001,
    'N_EPOCHS': 10,
    'CLASS_NAMES': ['sad', 'happy', 'angry', 'nothing']
}
```

```
import tensorflow as tf
import matplotlib.pyplot as plt
train_dataset = tf.keras.utils.image_dataset_from_directory(
    '/content/drive/MyDrive/EmotionsDataset_Splitted/data/test',
    image_size=(CONFIG['IM_SIZE'], CONFIG['IM_SIZE']),
    batch_size=32
)

CONFIG['CLASS_NAMES'] = train_dataset.class_names
CNN_INPUT_SIZE = cnn_model.input_shape[1:3]
print("CNN Model expects input shape:", CNN_INPUT_SIZE)
```

```
plt.figure(figsize=(12, 12))
for images, labels in train_dataset.take(1):
    images = images / 255.0
    for i in range(16):
        ax = plt.subplot(4, 4, i + 1)

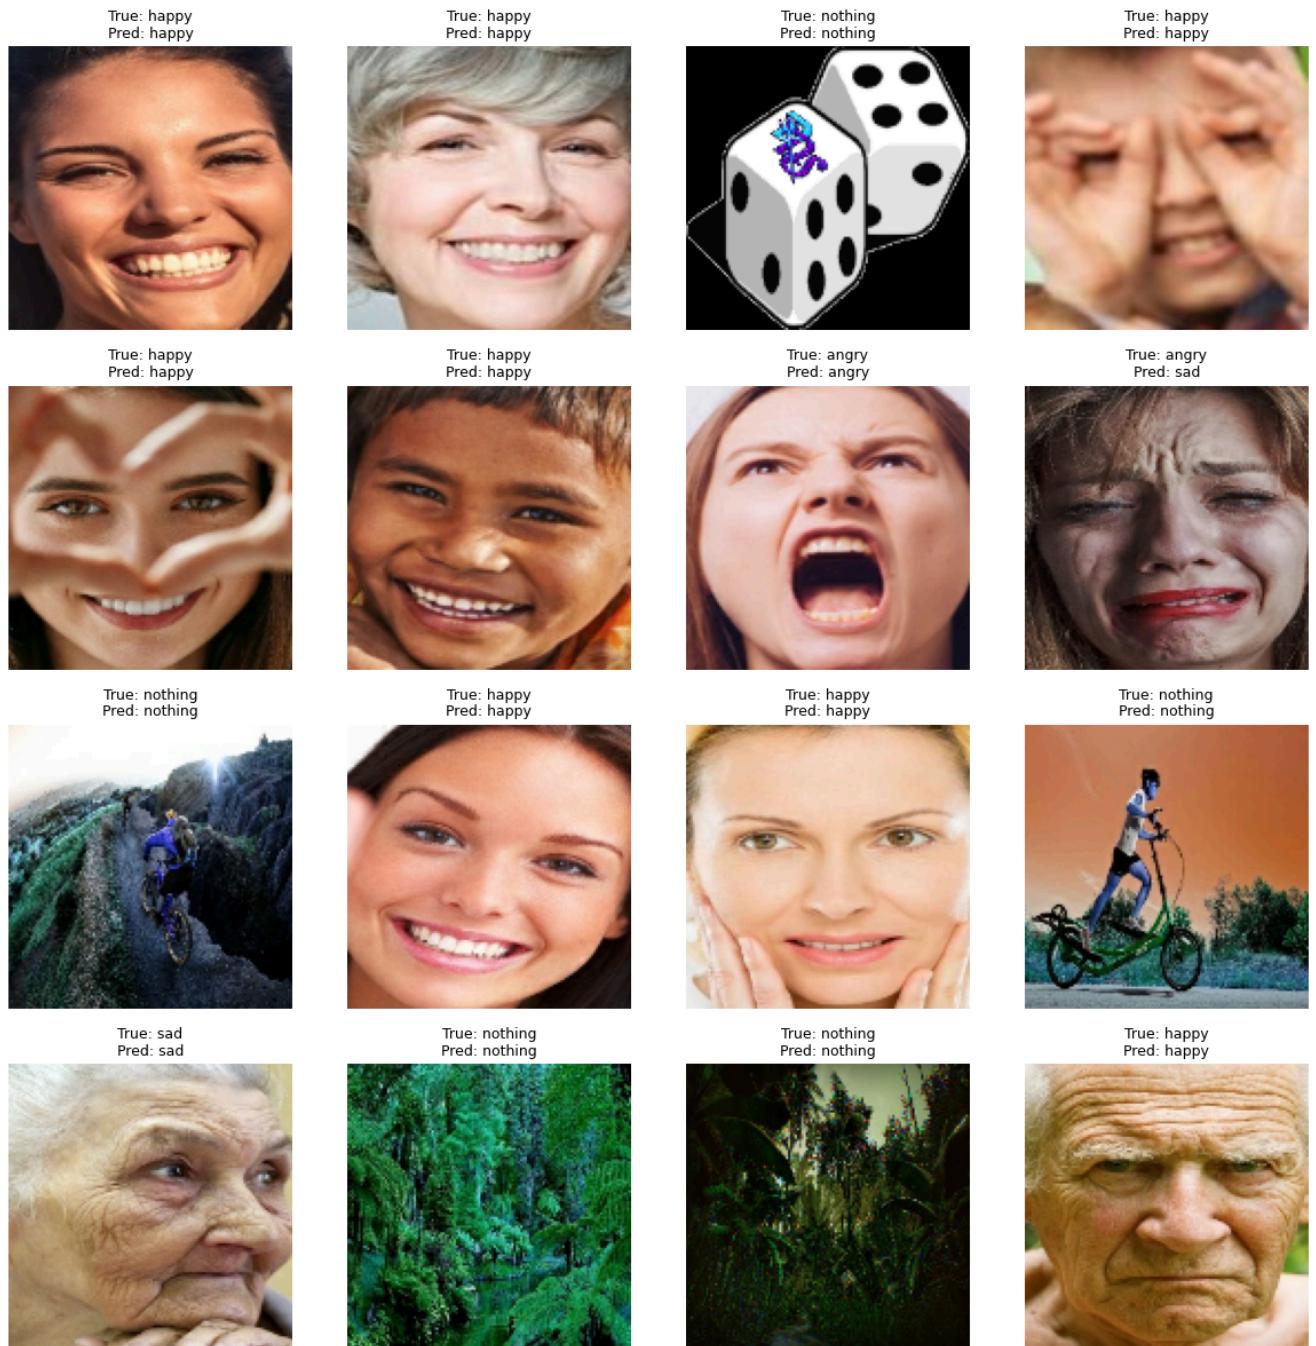
        img = tf.image.resize(images[i], CNN_INPUT_SIZE)
        img_expanded = tf.expand_dims(img, axis=0)
        preds = cnn_model(img_expanded, training=False)
        pred_label_idx = tf.argmax(preds, axis=1).numpy()[0]
        pred_label = CONFIG["CLASS NAMES"][pred_label_idx]

        true_label_idx = labels[i].numpy()
        true_label = CONFIG["CLASS NAMES"][true_label_idx]

        plt.imshow(img)
        plt.title(f"True: {true_label}\nPred: {pred_label}", fontsize=9)
        plt.axis("off")

plt.tight_layout()
plt.show()
```

Found 788 files belonging to 4 classes.  
CNN Model expects input shape: (128, 128)



## ▼ BERT MODEL

```
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Input, Conv2D, BatchNormalization, LeakyReLU, MaxPooling2D, Flatten, Dense, Dropout

CONFIG = {
    'IM_SIZE': 128,
    'DROPOUT_RATE': 0.5,
    'NUM_CLASSES': 10,
    'LEARNING_RATE': 0.001,
    'N_EPOCHS': 20
}
```

```
def darknet_block(filters, kernel_size, strides=1, padding='same'):
    return [
        Conv2D(filters, kernel_size, strides=strides, padding=padding, use_bias=False),
        BatchNormalization(),
        LeakyReLU(alpha=0.1)
    ]

def create_darknet_model():
    model = Sequential()
    model.add(Input(shape=(CONFIG['IM_SIZE'], CONFIG['IM_SIZE'], 3)))

    # Darknet-like Conv blocks
    for block in [
        (32, 3),
        (64, 3),
        (128, 3),
        (256, 3),
        (512, 3)
    ]:
        filters, k = block
        for layer in darknet_block(filters, k):
            model.add(layer)
    model.add(MaxPooling2D(pool_size=(2, 2)))

    # Classification head
    model.add(Flatten())
    model.add(Dense(256, activation='relu'))
    model.add(Dropout(CONFIG['DROPOUT_RATE']))
    model.add(Dense(CONFIG['NUM_CLASSES'], activation='softmax'))

    return model

darknet_model = create_darknet_model()
darknet_model.compile(
    optimizer=tf.keras.optimizers.Adam(learning_rate=CONFIG['LEARNING_RATE']),
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy']
)

darknet_model.summary()
```

```
↳ /usr/local/lib/python3.11/dist-packages/keras/src/layers/activations/leaky_relu.py:41: UserWarning: Argument `alpha` is deprecated.
  warnings.warn(
Model: "sequential_4"
```

Layer (type)	Output Shape	Param #
conv2d_13 (Conv2D)	(None, 128, 128, 32)	864
batch_normalization_10 (BatchNormalization)	(None, 128, 128, 32)	128
leaky_re_lu_10 (LeakyReLU)	(None, 128, 128, 32)	0
max_pooling2d_13 (MaxPooling2D)	(None, 64, 64, 32)	0
conv2d_14 (Conv2D)	(None, 64, 64, 64)	18,432
batch_normalization_11 (BatchNormalization)	(None, 64, 64, 64)	256
leaky_re_lu_11 (LeakyReLU)	(None, 64, 64, 64)	0
max_pooling2d_14 (MaxPooling2D)	(None, 32, 32, 64)	0
conv2d_15 (Conv2D)	(None, 32, 32, 128)	73,728
batch_normalization_12 (BatchNormalization)	(None, 32, 32, 128)	512
leaky_re_lu_12 (LeakyReLU)	(None, 32, 32, 128)	0
max_pooling2d_15 (MaxPooling2D)	(None, 16, 16, 128)	0
conv2d_16 (Conv2D)	(None, 16, 16, 256)	294,912
batch_normalization_13 (BatchNormalization)	(None, 16, 16, 256)	1,024
leaky_re_lu_13 (LeakyReLU)	(None, 16, 16, 256)	0
max_pooling2d_16 (MaxPooling2D)	(None, 8, 8, 256)	0
conv2d_17 (Conv2D)	(None, 8, 8, 512)	1,179,648
batch_normalization_14 (BatchNormalization)	(None, 8, 8, 512)	2,048
leaky_re_lu_14 (LeakyReLU)	(None, 8, 8, 512)	0
max_pooling2d_17 (MaxPooling2D)	(None, 4, 4, 512)	0
flatten_3 (Flatten)	(None, 8192)	0
dense_6 (Dense)	(None, 256)	2,097,408
dropout_3 (Dropout)	(None, 256)	0
dense_7 (Dense)	(None, 10)	2,570

```
import numpy as np

# Example: 100 RGB images of size 128x128 and their labels
X_train = np.random.rand(100, CONFIG['IM_SIZE'], CONFIG['IM_SIZE'], 3).astype(np.float32)
y_train = np.random.randint(0, CONFIG['NUM_CLASSES'], size=(100,))

train_data = tf.data.Dataset.from_tensor_slices((X_train, y_train)).batch(32)

darknet_model.compile(
    optimizer=tf.keras.optimizers.Adam(learning_rate=CONFIG['LEARNING_RATE']),
    loss=tf.keras.losses.SparseCategoricalCrossentropy(),
    metrics=[
        tf.keras.metrics.SparseCategoricalAccuracy(name='accuracy'),
        tf.keras.metrics.SparseTopKCategoricalAccuracy(k=2, name='top_k_accuracy')
    ]
)

history_darknet = darknet_model.fit(
    train_data,
    epochs=CONFIG['N_EPOCHS'],
    verbose=1
)
```

```
↳ Epoch 1/20
4/4 ━━━━━━━━━━━━━━━━ 12s 1s/step - accuracy: 0.0798 - loss: 11.2750 - top_k_accuracy: 0.1614
```

```

Epoch 2/20
4/4 ━━━━━━━━ 8s 2s/step - accuracy: 0.1793 - loss: 14.7179 - top_k_accuracy: 0.2752
Epoch 3/20
4/4 ━━━━━━ 6s 1s/step - accuracy: 0.1873 - loss: 7.5774 - top_k_accuracy: 0.3289
Epoch 4/20
4/4 ━━━━━━ 8s 2s/step - accuracy: 0.2736 - loss: 4.3722 - top_k_accuracy: 0.3902
Epoch 5/20
4/4 ━━━━━━ 6s 1s/step - accuracy: 0.2247 - loss: 3.0244 - top_k_accuracy: 0.3600
Epoch 6/20
4/4 ━━━━━━ 8s 2s/step - accuracy: 0.1480 - loss: 2.6373 - top_k_accuracy: 0.2370
Epoch 7/20
4/4 ━━━━━━ 7s 2s/step - accuracy: 0.2834 - loss: 1.9917 - top_k_accuracy: 0.4802
Epoch 8/20
4/4 ━━━━━━ 8s 2s/step - accuracy: 0.2782 - loss: 1.8180 - top_k_accuracy: 0.5099
Epoch 9/20
4/4 ━━━━━━ 7s 2s/step - accuracy: 0.3550 - loss: 1.6828 - top_k_accuracy: 0.5579
Epoch 10/20
4/4 ━━━━━━ 8s 2s/step - accuracy: 0.5045 - loss: 1.5674 - top_k_accuracy: 0.6553
Epoch 11/20
4/4 ━━━━━━ 6s 1s/step - accuracy: 0.5141 - loss: 1.3654 - top_k_accuracy: 0.6862
Epoch 12/20
4/4 ━━━━━━ 10s 1s/step - accuracy: 0.5778 - loss: 1.1295 - top_k_accuracy: 0.7877
Epoch 13/20
4/4 ━━━━━━ 8s 2s/step - accuracy: 0.6420 - loss: 1.0896 - top_k_accuracy: 0.8205
Epoch 14/20
4/4 ━━━━━━ 6s 1s/step - accuracy: 0.7099 - loss: 0.9110 - top_k_accuracy: 0.8832
Epoch 15/20
4/4 ━━━━━━ 8s 2s/step - accuracy: 0.8278 - loss: 0.6879 - top_k_accuracy: 0.9355
Epoch 16/20
4/4 ━━━━━━ 7s 2s/step - accuracy: 0.8176 - loss: 0.7575 - top_k_accuracy: 0.9138
Epoch 17/20
4/4 ━━━━━━ 8s 2s/step - accuracy: 0.7385 - loss: 0.6539 - top_k_accuracy: 0.8956
Epoch 18/20
4/4 ━━━━━━ 9s 2s/step - accuracy: 0.7620 - loss: 0.7121 - top_k_accuracy: 0.8306
Epoch 19/20
4/4 ━━━━━━ 8s 1s/step - accuracy: 0.8560 - loss: 0.4980 - top_k_accuracy: 0.9449
Epoch 20/20
4/4 ━━━━━━ 7s 2s/step - accuracy: 0.8363 - loss: 0.4434 - top_k_accuracy: 0.9291

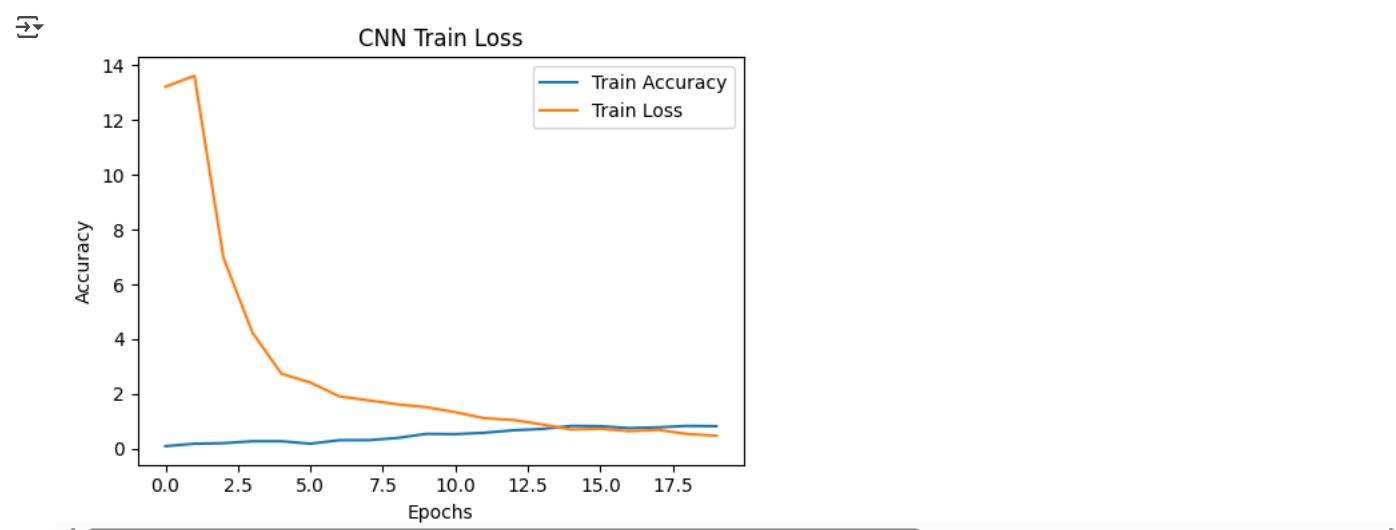
```

```

import matplotlib.pyplot as plt
plt.figure(figsize=(10, 4))
plt.subplot(1, 2, 1)
plt.plot(history_darknet.history['accuracy'], label='Train Accuracy')
plt.plot(history_darknet.history['loss'], label='Train Loss')
plt.title('CNN Train Accuracy')
plt.title('CNN Train Loss')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

plt.tight_layout()
plt.show()

```



```

darknet_eval_results = darknet_model.evaluate(train_data, verbose=1)
print(f"\nDarknet Model Evaluation Results:\nLoss: {darknet_eval_results[0]:.4f}, Accuracy: {darknet_eval_results[1]:.4f}, Top-2 Accuracy: {darknet_eval_results[2]:.4f}")

```

4/4 ━━━━━━ 3s 531ms/step - accuracy: 0.0928 - loss: 2.3114 - top\_k\_accuracy: 0.3202

Darknet Model Evaluation Results:  
Loss: 2.2877, Accuracy: 0.1200, Top-2 Accuracy: 0.3500

```
import cv2
import numpy as np
import matplotlib.pyplot as plt

img_path = '/content/drive/MyDrive/EmotionsDataset_Splitted/data/test/sad/sad_99.jpg'

test_image = cv2.imread(img_path)

if test_image is None:
    print("Image not found or failed to load:", img_path)
else:
    test_image = cv2.cvtColor(test_image, cv2.COLOR_BGR2RGB)
    test_image_resized = cv2.resize(test_image, (CONFIG['IM_SIZE'], CONFIG['IM_SIZE']))
    test_image_resized = test_image_resized.astype('float32') / 255.0
    test_image_resized = np.expand_dims(test_image_resized, axis=0)
    prediction = darknet_model.predict(test_image_resized)
    predicted_class = np.argmax(prediction, axis=1)
    plt.imshow(test_image)
    plt.axis('off')
    plt.show()
    print(f"Predicted class: {predicted_class[0]}")
```

⟳ 1/1 ━━━━━━ 0s 206ms/step



```
import cv2
import numpy as np
import matplotlib.pyplot as plt

img_path = '/content/drive/MyDrive/EmotionsDataset_Splitted/data/test/angry/angry_99.jpg'

test_image = cv2.imread(img_path)

if test_image is None:
    print("Image not found or failed to load:", img_path)
else:
    test_image = cv2.cvtColor(test_image, cv2.COLOR_BGR2RGB)
    test_image_resized = cv2.resize(test_image, (CONFIG['IM_SIZE'], CONFIG['IM_SIZE']))
    test_image_resized = test_image_resized.astype('float32') / 255.0
    test_image_resized = np.expand_dims(test_image_resized, axis=0)
    prediction = darknet_model.predict(test_image_resized)
    predicted_class = np.argmax(prediction, axis=1)
    plt.imshow(test_image)
    plt.axis('off')
    plt.show()
    print(f"Predicted class: {predicted_class[0]}")
```

1/1 0s 65ms/step



Predicted class: 8

```
import cv2
import numpy as np
import matplotlib.pyplot as plt

img_path = '/content/drive/MyDrive/EmotionsDataset_Splitted/data/test/happy/happy_99.jpg'

test_image = cv2.imread(img_path)

if test_image is None:
    print("Image not found or failed to load:", img_path)
else:
    test_image = cv2.cvtColor(test_image, cv2.COLOR_BGR2RGB)
    test_image_resized = cv2.resize(test_image, (CONFIG['IM_SIZE'], CONFIG['IM_SIZE']))
    test_image_resized = test_image_resized.astype('float32') / 255.0
    test_image_resized = np.expand_dims(test_image_resized, axis=0)
    prediction = darknet_model.predict(test_image_resized)
    predicted_class = np.argmax(prediction, axis=1)
    plt.imshow(test_image)
    plt.axis('off')
    plt.show()
    print(f"Predicted class: {predicted_class[0]}")
```

1/1 0s 64ms/step



Predicted class: 8

```
import cv2
import numpy as np
import matplotlib.pyplot as plt

img_path = '/content/drive/MyDrive/EmotionsDataset_Splitted/data/test/nothing/nothing_99.jpg'

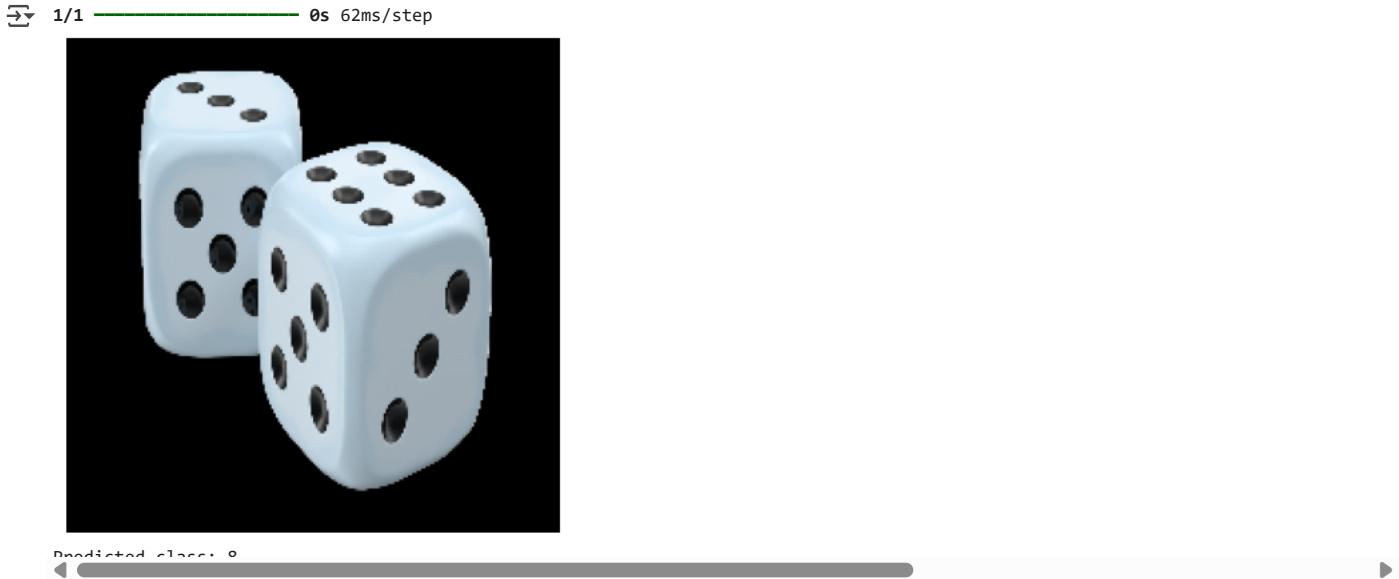
test_image = cv2.imread(img_path)

if test_image is None:
    print("Image not found or failed to load:", img_path)
else:
```

```

test_image = cv2.cvtColor(test_image, cv2.COLOR_BGR2RGB)
test_image_resized = cv2.resize(test_image, (CONFIG['IM_SIZE'], CONFIG['IM_SIZE']))
test_image_resized = test_image_resized.astype('float32') / 255.0
test_image_resized = np.expand_dims(test_image_resized, axis=0)
prediction = darknet_model.predict(test_image_resized)
predicted_class = np.argmax(prediction, axis=1)
plt.imshow(test_image)
plt.axis('off')
plt.show()
print(f"Predicted class: {predicted_class[0]}")

```



```

CONFIG = {
    'IM_SIZE': 128,
    'DROPOUT_RATE': 0.5,
    'NUM_CLASSES': 3,
    'LEARNING_RATE': 0.001,
    'N_EPOCHS': 10,
    'CLASS_NAMES': ['sad', 'happy', 'angry', 'nothing']
}

```

```

import tensorflow as tf
import matplotlib.pyplot as plt

train_dataset = tf.keras.utils.image_dataset_from_directory(
    '/content/drive/MyDrive/EmotionsDataset_Splitted/data/test',
    image_size=(CONFIG['IM_SIZE'], CONFIG['IM_SIZE']),
    batch_size=32
)

CONFIG['CLASS_NAMES'] = train_dataset.class_names
CNN_INPUT_SIZE = darknet_model.input_shape[1:3]

print("Darknet Model expects input shape:", CNN_INPUT_SIZE)
plt.figure(figsize=(12, 12))
for images, labels in train_dataset.take(1):
    images = images / 255.0
    for i in range(16):
        ax = plt.subplot(4, 4, i + 1)

        img = tf.image.resize(images[i], CNN_INPUT_SIZE)

        img_expanded = tf.expand_dims(img, axis=0)
        preds = cnn_model(img_expanded, training=False)
        pred_label_idx = tf.argmax(preds, axis=1).numpy()[0]
        pred_label = CONFIG["CLASS_NAMES"][pred_label_idx]
        true_label_idx = labels[i].numpy()
        true_label = CONFIG["CLASS_NAMES"][true_label_idx]

        plt.imshow(img)
        plt.title(f"True: {true_label}\nPred: {pred_label}", fontsize=9)
        plt.axis("off")

plt.tight_layout()
plt.show()

```